

Final Project

Machine Learning 2

Author: Lina Faik

Date: May 2025

Disclaimer: This presentation is for educational purposes only and should not be reused or redistributed without explicit permission.

Table of contents

01

Objectives

02

Project Ideas

03

Guidelines

04

Evaluation Criteria

05

Q&A

01

Objectives

01. Objectives

✓ **Apply Machine Learning Techniques:**

Utilize advanced machine learning algorithms learned in the course to solve real-world problems.

✓ **End-to-End Project Execution:**

Experience the entire machine learning workflow from data collection, preprocessing, and model training to evaluation and deployment.

✓ **Critical Thinking and Problem Solving:**

Develop the ability to identify, analyze, and solve complex problems using data-driven approaches.

✓ **Collaborative Work:**

Enhance teamwork and project management skills by working effectively in groups.

✓ **Communication Skill:**

Improve the ability to communicate technical results and insights through reports and presentations.

✓ **Practical Implementation:**

Gain hands-on experience with tools and libraries used in the industry, such as Python, scikit-learn, TensorFlow, or similar frameworks.

02

Project Ideas

02. Project Ideas

Overview

Team Formation:

- Form teams of 3-4 members.
- Individual projects are not allowed.

Project Selection:

- Choose a project from the provided list.
- You are also welcome to propose a different project topic if it aligns more with your interests.

Objectives and Data Sources:

- The listed objectives and data sources are suggestions.
 - You are responsible for defining your project's specific problem statement and validating the relevance of your data sources.
-

02. Project Ideas

Overview

Summary of the projects.

Provided as a guideline. Specifics will depend on your chosen approach and tools.

	Difficulty (1-5)	Topics
Project 1: Sentiment Analysis on Social Media	2/5	NLP
Project 2: Customer Segmentation	3/5	Clustering
Project 3: Recommendation System for E-commerce	4/5	NLP
Project 4: Stock Price Prediction	3/5	Time Series
Project 5: Predicting Customer Churn	2/5	Clustering, ML
Project 6: Predictive Maintenance	3/5	ML, Time Series
Project 7: Loan Default Prediction	2/5	ML
Project 8: Music Genre Classification	4/5	ML / DL
Project 9: Predicting Employee Attrition	2/5	ML
Project 10: Energy Consumption Forecasting	3/5	Time Series

02. Project Ideas

Project 1: Sentiment Analysis on Social Media

Context:

- Social media platforms are a rich source of public opinion.
- Understanding sentiment can help brands and organizations gauge public perception.
- Real-time sentiment analysis can inform marketing and communication strategies.

Goal:

- Develop a model to classify the sentiment of social media posts (positive, negative, neutral).
- Analyze trends in sentiment over time or across different topics.
- Visualize the results to provide actionable insights.

Data:

- [Kaggle Twitter US Airline Sentiment](#)
- [Sentiment140](#)
- [Social Media Sentiment Data](#)

Challenges:

- Preprocessing text data (removing noise, handling slang, emojis).
- Balancing the dataset if there is an uneven distribution of sentiment classes.
- Managing the large volume of data and ensuring the model scales effectively.

02. Project Ideas

Project 2: Customer Segmentation

Context:

- Businesses need to understand their customer base to tailor marketing strategies.
- Segmentation helps in identifying groups with similar behaviors and preferences.
- Data-driven segmentation can improve customer satisfaction and retention.

Goal:

- Use clustering techniques to segment customers based on purchasing behavior.
- Analyze the characteristics of each segment.
- Provide recommendations for targeted marketing campaigns.

Data:

- [Kaggle Customer Segmentation](#)
- [Mall Customer Segmentation Data](#)
- [Online Retail Dataset](#)

Challenges:

- Determining the optimal number of clusters.
 - Interpreting the clusters in a business context.
 - Ensuring the model scales with large datasets.
-

02. Project Ideas

Project 3: Recommendation System for E-commerce

Context:

- E-commerce platforms benefit from personalized recommendations to enhance user experience.
- Recommendations can increase sales and user engagement.
- Machine learning can automate and improve the recommendation process.

Goal:

- Build a recommendation system to suggest products to users based on their behavior.
- Compare collaborative filtering and content-based approaches.
- Evaluate the effectiveness of the recos using metrics like precision and recall.

Data:

- Kaggle Amazon Product Data
- MovieLens Dataset
- Goodreads Books

Challenges:

- Handling sparse data in user-item interaction matrices.
 - Scaling the recommendation system for large datasets.
 - Balancing diversity and relevance in recommendations.
-

02. Project Ideas

Project 4: Stock Price Prediction

Context:

- Financial markets are influenced by numerous factors, making prediction challenging.
- Accurate predictions can guide investment strategies.
- Historical data can be analyzed to identify patterns and trends.

Goal:

- Develop a model to predict future stock prices based on historical data.
- Explore different time series forecasting methods.
- Evaluate the model's performance and identify potential risks.

Data:

- [Yahoo Finance Historical Data](#)

Challenges:

- Handling noisy and non-stationary data.
 - Incorporating external factors like news and economic indicators.
 - Preventing overfitting to historical trends.
-

02. Project Ideas

Project 5: Predicting Customer Churn

Context:

- Customer retention is crucial for business profitability.
- Identifying factors leading to churn can help in devising retention strategies.
- Machine learning can automate the identification of at-risk customers.

Goal:

- Build a model to predict whether a customer will churn based on their behavior.
- Analyze the features that contribute most to churn.
- Provide recommendations for reducing churn.

Data:

- [Kaggle Telco Customer Churn](#)
- [Customer Churn Data](#)

Challenges:

- Balancing the dataset if there is an uneven distribution of churned vs. retained customers.
 - Feature engineering to capture relevant customer behaviors.
 - Interpreting the model to provide actionable business insights.
-

02. Project Ideas

Project 6: Predictive Maintenance

Context:

- Equipment failures can lead to costly downtime and repairs.
- Predictive maintenance aims to predict failures before they occur.
- Machine learning can analyze sensor data to identify early signs of failure.

Goal:

- Develop a model to predict equipment failures based on sensor data.
- Identify the most critical factors contributing to failures.
- Provide recommendations for maintenance schedules.

Data:

- [NASA Turbofan Engine Degradation Simulation Data Set](#)
- [Predictive Maintenance Dataset](#)
- [Predictive Maintenance Dataset 2](#)

Challenges:

- Handling large volumes of high-dimensional sensor data.
 - Balancing the dataset if failures are rare.
 - Interpreting the model to provide actionable maintenance insights.
-

02. Project Ideas

Project 7: Loan Default Prediction

Context:

- Financial institutions need to assess the risk of loan defaults.
- Accurate predictions can help in making informed lending decisions.
- Machine learning can analyze historical data to predict defaults.

Goal:

- Build a model to predict whether a loan will default based on applicant data.
- Identify the most significant factors contributing to defaults.
- Provide recommendations for risk mitigation.

Data:

- Kaggle Lending Club Loan Data
- Loan Default Data
- Home Credit Default Risk

Challenges:

- Handling imbalanced datasets where defaults are less frequent.
 - Feature engineering to capture relevant applicant characteristics.
 - Interpreting the model to provide actionable insights.
-

02. Project Ideas

Project 8: Music Genre Classification

Context:

- Music streaming services can benefit from automatic genre classification.
- It enhances user experience by enabling better music recommendations.
- Audio features can be analyzed to classify genres.

Goal:

- Develop a model to classify music tracks into genres based on audio features.
- Explore different feature extraction techniques.
- Evaluate the model's performance and interpret the results.

Data:

- [GTZAN Music Genre Dataset](#)
- [FMA: A Dataset For Music Analysis](#)
- [Million Song Dataset](#)

Challenges:

- Extracting relevant features from audio data.
 - Handling class imbalances if certain genres are less represented.
 - Ensuring the model generalizes well to new and diverse music tracks.
-

02. Project Ideas

Project 9: Predicting Employee Attrition

Context:

- Employee attrition can impact organizational productivity and morale.
- Predicting attrition can help in retaining valuable employees.
- Machine learning can analyze employee data to identify attrition risks.

Goal:

- Build a model to predict whether an employee will leave the company.
- Analyze the factors contributing most to attrition.
- Provide recommendations for improving employee retention.

Data:

- IBM HR Analytics Employee Attrition & Performance
- Employee Attrition Data
- HR Analytics Dataset

Challenges:

- Handling imbalanced datasets if attrition is less frequent.
 - Feature engineering to capture relevant employee characteristics.
 - Interpreting the model to provide actionable insights for HR practices.
-

02. Project Ideas

Project 10: Energy Consumption Forecasting

Context:

- Accurate energy consumption forecasting can help in managing resources.
- It is crucial for energy providers to balance supply and demand.
- Machine learning can analyze historical data for accurate forecasting.

Goal:

- Develop a model to forecast future energy consumption based on historical data.
- Identify patterns and trends in energy usage.
- Evaluate the model's performance and provide recommendations.

Data:

- [UCI Energy Consumption Dataset](#)
- [UK National Grid Data](#)
- [Household Electricity Consumption](#)

Challenges:

- Handling Seasonality and Trends: Capturing daily and seasonal consumption patterns.
 - Incorporating External Factors: Integrating weather, holidays, and economic activity data.
 - Dealing with Data Quality: Managing missing values and anomalies in historical data.
-

03

Guidelines

03. Guidelines

Project Planning:

- Define clear objectives and deliverables.
- Create a project timeline with milestones.

Data Handling:

- Identify and validate data sources.
- Clean and preprocess data, handling missing values.

Model Development:

- Choose suitable machine learning algorithms.
- Perform feature engineering and model training.

Evaluation:

- Use appropriate evaluation metrics.
- Implement cross-validation to ensure robustness.

Insights:

- Understand and explain model predictions.
- Derive practical insights and recommendations.

Collaboration:

- Assign team roles and responsibilities.
 - Hold regular team meetings and maintain documentation.
-

03. Guidelines

Presentation:

- Create clear visualizations.
- Prepare a final report and presentation.

Resources:

- Use tools like Python, scikit-learn, TensorFlow.
- Refer to course materials and online resources.

Feedback:

- Seek peer and instructor feedback.
 - Iterate on your model and approach as needed.
-

04

Evaluation Criteria

04. Evaluation Criteria

1. Project Planning and Execution (10%):

- Clarity and feasibility of project objectives.
- Adherence to the project timeline and milestones.
- Effective project management and teamwork.

2. Data Handling (20%):

- Quality and relevance of the data sources used.
- Effectiveness of data preprocessing and cleaning.
- Appropriateness of feature engineering.

3. Model Development (20%):

- Selection and justification of machine learning algorithms.
- Model training and tuning processes.
- Comparison and performance of different models.

4. Evaluation and Validation (20%):

- Use of appropriate evaluation metrics.
 - Implementation of cross-validation techniques.
 - Robustness and generalization of the model.
-

04. Evaluation Criteria

5. Insights and Interpretation (10%):

- Clarity in explaining model predictions.
- Derivation of actionable insights and recommendations.
- Understanding of feature importance and model behavior.

6. Presentation and Reporting (10%):

- Quality and clarity of visualizations.
- Comprehensive and well-structured final report.
- Effectiveness of the final presentation.

7. Ethical Considerations (5%):

- Addressing potential biases in data and models.

8. Innovation and Creativity (5%):

- Originality and creativity in problem-solving.
- Implementation of innovative techniques or approaches.

Q&A

Thank you!
